

Salience, Ranking, and Metabolism: Three Conflated Signals in Long-Running Agent Memory Systems

Baofeng Zhao

Independent Researcher

Abstract

Long-running agent memory systems commonly treat "importance" as a single quantity. We argue that it conflates three distinct signals, and that the conflation is a structural design error rather than an implementation detail. A scoring formula written in 2023 for a simulation demo (relevance plus recency plus importance) moved into production systems unexamined, letting the three jobs of "importance" (sedimentation, ranking, retirement, later separated here as salience, ranking, and metabolism) interfere with one another within a single signal. Our position rests on a service chain: memory serves the LLM; the LLM serves the human. And memory must be *memory*: self-describing and self-metabolizing, its salience and retirement decided by data and lifecycle rather than one-shot human-delegated verdicts (Invariant Io). Skip the middle link and even the most elaborate memory system drifts toward RAG. Since March 2026, over about seven months of production telemetry on a single conversational system (~38K raw memory items, ~148K association links as of the 2026-08 audit, the live store now holding ~108K), we argue that "importance" is at least three orthogonal signals that are best kept physically separated: salience (intrinsic to the data), ranking (a query-time service), and metabolism (lifecycle retirement). We document three production failures of conflation: a 93× semantic skew inside one counter, a structure-blind quality metric reporting 0.999 while the store's structure-aware coherence read 0.101, and 88% of raw memories being "islands" (73% at the time of writing, still declining without cleanup), yet containing the most structurally central roots. The retrieval-side pressure is measurable: an access-distribution Gini coefficient of 0.960, with the top 1% of nodes absorbing 51.9% of all accesses. We contribute a signal taxonomy with four invariants (Io–I3), seven design laws, and two same-day audit tools (the Silence Test and the RAG Test). No benchmark supremacy is claimed; all evidence comes from our own failures, of which we have plenty.

Keywords: agent memory; signal separation; positive-feedback bias; RAG; memory governance; position paper

1. Introduction: For Whom Does Memory Serve—the Legacy of RAG

"Can thinking take effect next time?—But you've already forgotten what you sorted out last time."

—an utterance from our system's corpus (lightly redacted)

Before any scoring formula, a more fundamental question: for whom does memory exist?

Human memory is not built for queries: an event sinks, and years later surfaces unprompted, as a scent, a tone of voice, a habit of handling things. Humans forget by *living*—not by executing a deletion policy, but by being slowly overgrown with later life. Human memory has no index, and it never grades its own recollections.

An LLM starts from the opposite direction: by default every conversation begins fresh. Its memory cannot grow organically; it must be built.

Our position is a service chain: memory serves the LLM; the LLM serves the human. Remove the middle link (let memory serve the human's queries directly), and no matter how elaborately the system is dressed up (extraction, consolidation, graphs, layers), we contend it functions, in essence, as RAG [10]: a person queries, a machine searches, the results are pasted into context with the model as glue. RAG as a tool is beyond reproach; as a memory architecture, we hold that it is a category error: it compresses the entire life of memory into the moment of retrieval, shaped by the query. Memory becomes an index; the model becomes a pipe.

Human memory and silicon memory are two different media; between media you need a bridge, not direct wiring. Pasting human documents raw into context is wiring; letting raw utterances settle into structure, letting derivations converge around anchors, keeping causal chains traceable, rebuilding the past in a way the model can grow on: that is a bridge. Our stance toward human intervention follows the same line: the structure of memory should be *recognized*, not manufactured. A bridge that can stand needs at least three different materials, which are our protagonists.

The service chain answers *for whom* memory serves; a second, equally fundamental question remains: *by what right* does memory hold together? Our secondary thesis: to give an LLM memory, the precondition is that memory be memory itself, not memory under human interference. Saliency is not a score someone assigns; it is the data's own shape. Metabolism is not a deletion policy someone executes; it is memory living through its own lifecycle. The semantics of a database are binary: a record is either present or absent, and its verbs are insert, update, delete, query; the semantics of memory are alive (grow, converge, retire, circulate), all decided endogenously by the data: data decides saliency; data decides metabolism. Asking an LLM for a 1–10 score at write time, or an LLM judge for keep-or-delete verdicts, inserts a foreign verdict between the data and its own

weight: the signature of human interference. A concurrent work [7] asks, "Is agent memory a database?" Our answer: no, and we offer the positive program of what it is: a self-consistent, metabolizing memory organism.

The current industry default is a retrieval-age one: it treats memory as something to be optimized for lookup. The concrete fossil of the retrieval mentality is a formula inherited from a demo. In 2023, a simulation of a town of 25 LLM agents (Generative Agents [1]) proposed what is now effectively the industry-default memory retrieval formula:

score = relevance (cosine similarity) + recency (exponential decay, factor 0.995, computed from last access) + importance (an LLM-assigned 1–10 score at write time)

The formula was written for a performance, yet became the skeleton of many production systems. The problem is not the summation (weighted fusion is innocent) but that it quietly lets one signal work three jobs: how *thick* a memory is (whether it is an anchor that later content converges around), where it *rank*s this time (whether it fits the present query), and how much longer it *lives* (cleanup logic widely reads counter/popularity-style inputs). Stuffing three signals into one field is like engraving heart rate, speed, and fuel consumption onto a single dial: the needle moves, but it no longer reports any one thing.

Our claim: in the system we studied, this is not roughness of implementation but a structural category error. Four contributions, each independently checkable:

- **C1 (Thesis)** A repositioning of the memory service chain: memory serves the LLM, the LLM serves the human; memory serving human queries directly tends to collapse into RAG.
- **C2 (Taxonomy)** A formal separation of the salience/ranking/metabolism signals, the memory-self-consistency principle, and four invariants (I0–I3).
- **C3 (Evidence)** Three production failure records (a 93× field-semantics skew, a 0.999 false-health metric, 88% island roots) with post-separation field readings (93× → 2.2×; Gini 0.960).
- **C4 (Tools)** Two same-day executable audit tools: the Silence Test and the RAG Test.

Roadmap: §2 locates the lesion in related work → §3 gives the diagnostic framework → §4 submits the case files → §5–§6 derive design laws and audit tools → §7–§8 bound the claims and close. Finally, three readings to forestall. This paper is not against retrieval; retrieval is memory's legitimate act of access; we object only to its doubling as sole arbiter. It is not against scoring; a score may serve as input to the ranking layer, not as a proxy for salience or a termination criterion for metabolism. It submits no new architecture; what we deliver is a checkup framework and invariants that any existing system can self-test against without rebuilding.

2. Related Work: Three Heirs and a Critical Wave

(a) The textbook form—retrieval as reinforcement. In Generative Agents [1], recency decays from last access: a retrieved memory has its recency refreshed and thus becomes easier to retrieve again. This is the well-studied positive-feedback/popularity bias of recommender systems [5]; there its endpoint is called a filter bubble. Moved into a memory system, the endpoint is a memory bubble: a model left with only the things it likes to recall. To our knowledge, this specific critique has not been formalized for the agent-memory context.

(b) The fused form—one judge, three jobs. Memo [2] makes two LLM calls per write: extract candidate facts, then an LLM judge rules ADD / UPDATE / DELETE over vector-similar neighbors. Saliency judgment (how important is the new fact), metabolism decision (should the old fact die), and retention material (what survives for retrieval) are fused in a single ruling, triggered by the accident of "whoever similarity can retrieve." In fairness, the field is starting to feel the pain: the Memo team now discusses eviction and forgetting [3], and the community proposes access-frequency-driven lifecycle cleanup [4]. Where such work begins to separate retrieval decay from eviction, our distinct contribution is to generalize it into a full three-signal taxonomy with invariants (I0–I3) and to supply longitudinal production evidence for each failure mode; the problem is acknowledged, but a remedy is not yet formulated.

(c) The half-separated form—closest, one step short. Zep's [6] bi-temporal edges let superseded facts be invalidated rather than deleted (history preserved), the most lucid handling of metabolism. But its retrieval ranking remains relevance + recency + episode usage; usage frequency enters ranking directly; behavioral contamination of the measure persists. MemGPT [12] popularized tiered memory management (paging between a working context and external storage), but its eviction likewise reads recency and usage.

(d) The gathering critical wave. The concurrent work "Is Agent Memory a Database?" [7] criticizes current scoring as "local heuristics applied at retrieval time." We agree with the diagnosis and move two steps further: a complete three-signal taxonomy with invariants, and production-grade longitudinal failure evidence. [7] also proposes GEM, a framework whose correctness is defined at the state-trajectory level (four state-level operators—ingestion, revision, forgetting, retrieval—plus six correctness properties); our invariants instead define correctness at the signal level (fields and their writers). The two are complementary rather than competing: GEM governs what a memory operation must do to a state, whereas we govern which writer may touch which field. A wave of criticism is forming; this paper contributes a frame for it.

(e) An anticipated rebuttal: why not just ask an LLM to score? The most natural defense: let an LLM score importance at write time ([1]'s 1–10), or judge at update time ([2]), does that not give us "importance"? Three reasons it fails. First, it violates I0: a one-shot external verdict substitutes for the data's own weight, and the judge cannot see the future; on write day no one knows which raw utterance will become the sole basis of a

critical judgment three months later (§4.3). Second, the judge is retrieval's prisoner: an LLM judge only sees the neighbors similarity can put in front of it, so "what counts as important" is hijacked by "what is similar" (§2b). Third, verdicts cannot metabolize: once written, a score becomes a lifetime label, and once written, the score becomes a permanent label, effectively a one-time in/out verdict.

Table 2: Three heirs × handling of the three signals

System	Salience (intrinsic)	Ranking (query service)	Metabolism (lifecycle)
Generative Agents [1]	none (LLM 1–10 at write, for life)	blended relevance + recency (access-refreshed) + importance	none
Memo [2]	fused with metabolism & retention in one LLM ruling	vector similarity + ruling outcome	DELETE inside the same ruling (triggered by similar neighbors)
Zep [6]	none	relevance + recency + usage (behavior in ranking)	bi-temporal invalidate-not-delete (closest to I2)
This paper	intrinsic to data, convergence measured in-scope, unaffected by behavior	bounded [0,100] + decay, never written back	item-level cascade, lifecycle-triggered, popularity-blind

3. Method: Signal Definitions, Invariants, and a Feedback Analysis

Notation: $\sigma(x)$ is the salience of item x , $\rho(x)$ its ranking-layer score, $\mu(x)$ its metabolic state; the three have distinct domains and distinct writers; this is the section's entire claim. On ρ : the stored score is an input to per-query ordering, fused at query time with query similarity; the persistent score itself carries no query.

Definition 1 (Salience). A structural property of x , measured relative to the store: *how much derived content converges around x as anchor*, measured within the correct scope (counting across semantic domains yields fan-out noise, not convergence). Key property: unaffected by retrieval behavior. It is the shape of the data, not its popularity; it grows while no one retrieves. Concretely, our implementation measures it as a normalized convergence density: $\sigma(x) = \log_{1p}(\text{physical in-degree from anchor-forming edges}) / \log_{1p}(p_{99})$, where p_{99} is the 99th-percentile in-degree over those same anchor edges, counted only over anchor edge types (raw_basis, ai_basis, derived_from); recall-driven edges are excluded by construction, so σ stays flat under retrieval. One honest caveat: the p_{99} denominator is a store-level statistic, so σ is a within-store relative measure; a node's σ

can drift as the store grows even if its own in-links do not. σ compares items within one store and epoch, not across stores.

Definition 2 (Ranking). The output of a query-time algorithmic service; a bounded counter (e.g., $[0,100]$) with forgetting-curve decay is a legal implementation. Key properties: recomputable at will; serves only the consumer; never written back to data. Our implementation: an Ebbinghaus-style score (following MemoryBank [11]) $\rho = \text{count_norm} \times \exp(-\Delta t / S)$, where count is the recall count (hard-capped at 100), count_norm is that count normalized to $[0,1]$ by the cap, Δt is days since last recall, and $S = 7 \times (1 + \text{count})$ carries the same day units as Δt .

Definition 3 (Metabolism). A lifecycle decision, cascading at the item level (when a raw utterance retires, its derivations and links retire with it). Key property: triggered by lifecycle events, not by any popularity or ranking score. Low frequency $\neq$ expired—coldness is the birthright of root nodes, not their crime. Our implementation: retirement is a soft status mark with item-level cascade; it reads lifecycle events only, never a counter. Figure 1 summarizes the three domains and the walls that keep them apart.


A deeper footnote to the separation. Ranking's recency decay treats time as a unit of measurement; nothing wrong there: the clock is the ranking layer's legitimate tool. But salience is not a function of the clock; it is sediment of interaction, constituted by how much derivation converges, regardless of the wall clock. Time belongs to the clock; interaction belongs to memory. This is the seam along which ranking and salience must part.


Invariants:

- **I0 (No proxy):** memory's persistent properties (salience, survival) must not be assigned by one-shot external verdicts (LLM scores, LLM judge rulings); they are endogenous to data and lifecycle;
- **I1 (No contamination):** no change in ranking signals may alter salience or the data itself;
- **I2 (No wrongful death):** metabolism must not be triggered solely by ranking/popularity scores;
- **I3 (No polysemy):** one physical field must not carry multiple semantics.

Feedback sketch (why conflation necessarily worsens) (Figure 3). Let ranking score s depend on recency r ; r decays over time but refreshes on access; then $\partial s / \partial \text{access} > 0$, while access probability rises with s —a self-reinforcing loop. Only by locking "access" strictly inside the ranking layer and forbidding any write-back to persistent fields (I1) is the loop physically cut. This is the minimal structural requirement for cutting a feedback loop, not an aesthetic preference. The magnitude of the pressure is measurable: in this system's current (post-separation, bounded-and-capped) ranking layer: 25,705 nodes, an access-count Gini coefficient of 0.960 (Figure 7), the top 1% of nodes absorbing 51.9% of all accesses, the top 10% taking 96.9%, and 85.8% of nodes with zero accesses. The correct reading: winner-take-all in the ranking layer is legal; it is confined to the service layer,

recomputable at will. The point of the number is the warning: a distribution like this, written into any persistent field, is a retirement sentence for 85.8% of zero-access nodes, and most of them are roots.





Written into any persistent field, this distribution is a retirement sentence
for the 85.8% zero-access nodes (most of them roots) — this is why I1 exists

Figure 7 | Lorenz curve of the ranking-layer access distribution (production-measured, 2026-08). 25,705 nodes: Gini 0.960; the top 1% of nodes take 51.9% of all accesses; the top 10% take 96.9%; 85.8% have zero accesses. Measured after the three-signal separation — winner-take-all in the ranking layer is legal (bounded, recomputable, never written back); the figure quantifies the pressure and explains why no persistent field may reference such a distribution (Invariant I1).

The life of one memory (a worked example) (Figure 2). Put the three definitions on one timeline. Day 0: a raw utterance enters; in the wired world it receives an LLM score of 5 and lives with that lifetime label; in the bridged world it has no score, only an identity—raw anchor. Day 30: it happens to be retrieved twice; in the wired world its recency refreshes, its score rises, and it begins to attract more retrieval; in the bridged world its ranking rises too, but all of it stays inside the service layer; its salience is untouched, still zero convergence. Day 90: no one asks; in the wired world its score decays to zero and it enters the cleanup queue; in the bridged world it is equally "cold," but only its ranking is cold, not its right to exist. Day 120: a new judgment needs a basis, and finds this is the only anchor; in the wired world it was deleted by popularity on day 100; the judgment loses its basis while the dashboard glows green; in the bridged world it is surrounded by derivation, and its salience grows for the first time, from island to root. One memory, two fates; the difference lies not in any single algorithmic step, but in whether the three jobs are separated.

The human-memory contrast (and one honest boundary). Human retrieval does strengthen memory; the testing effect is a classic result of cognitive psychology [8]; is I1's "no access write-back" anti-biological? No—biology is being quoted halfway: human strengthening occurs in a separate synaptic substrate, bounded by homeostatic renormalization, and rebalanced wholesale during sleep consolidation [9]: reinforcement with a

ceiling, a ledger, and a rebalancing. Software counters have only reinforcement, no homeostasis: the unbounded positive feedback is exactly the memory bubble of §2a. I1's correct reading is not "don't imitate biology" but "imitate all of it": access sensitivity may live in the ranking layer (bounded, decaying, recomputable); the persistent layer must remain behavior-independent, unless a separate mechanism with homeostasis and rebalancing is introduced.




Figure 1 | Three-signal architecture. Data domain, service domain, and lifecycle domain are physically separate. Saliency σ lives in the data domain (structural, behavior-independent); ranking ρ in the service domain (bounded + decay, recomputable, never written back — Wall I1); metabolism μ in the lifecycle domain (item-cascade, popularity barred — Wall I2); I3 bars polysemous fields; I0 bars external verdicts on persistent properties. The dashed red cross marks the access write-back path cut by I1.

4. Case Study: Failure Records from a Single Production System (n=1)

System background: a conversational agent memory system running continuously since March 2026 (~7 months), dual-store (raw utterance layer + distillation layer), ~38,000 memory items, ~148,000 association links (2026-08 audit snapshot; the live store now holds ~108K). All numbers below come from real replays of the production store.

Table 1: Evidence timeline (single store).

time point	what happened	cohort
early audit (pre-fix)	full-store audit finds the 93× field-semantics skew and 88% island share	whole raw layer at audit

time point	what happened	cohort
post-fix (field split, cap [0,100])	cross-type span $93\times \rightarrow 2.2\times$; peak pinned at 100	same field, re-measured
2026-08 (fixed snapshot)	island share 67.8% (user, n=5,956) / 82.6% (AI, n=3,173); combined 73.0% (n=9,129); Gini 0.960 (n=25,705)	raw utterances that month
2026-10 (live store)	~108K links; raw layer ~13,364 (user 6,289 + AI 7,075)	current


Glossary (terms as used in this paper): *cortex* — the distilled-node store (~38k nodes); *raw layer* — verbatim utterances; *distillation layer* — derived nodes; *anchor* — a node that derived content converges on; *standing-wave pool* — live cortex nodes with completed vectors above the activity gate; *activity gate* — the weight threshold admitting a node to that pool.

How to read this section: these are identification-and-handling records from one system, not controlled trials. Splitting a field shared by three semantics necessarily compresses the cross-type span ($93\times \rightarrow 2.2\times$ is partly the definitional consequence of the split), and the peak sits at 100 because the cap is 100. The island-share change ($88\% \rightarrow 73\%$) is an observation over time without a control group; causal attribution is not established.

4.1 First page: one field, three semantics, a silent $93\times$ distortion

In the early design, one count field was written by three paths: retrieval +1, distillation +1, another subsystem +1. No one noticed, until an early full-store audit produced the X-ray: "principle"-type items averaged 449, "raw utterance"-type items averaged 4.8. Same field, same decay curve, a $93\times$ span across semantics; a single peak of 1965 became a black-hole item absorbing all ranking attention. We note that the span is a ratio of type means and is dominated by that single item; the pre-split per-item distribution is no longer reconstructible (the fix itself split the field), so medians cannot be reported post hoc. Every ranking and cleanup routine reading that field was doing arithmetic on three incommensurable quantities daily, while the reports stayed perfectly normal.


The fix was not tuning (tuning cannot save a field shared by three semantics) but splitting: one field per semantics, one decay constant each, uniformly capped at [0,100]. Post-separation field data (measured 2026-08): the cross-type span compressed from $93\times$ to $2.2\times$ (principle 2.9 vs raw 1.3; the residual gap is legitimate under unified semantics); the peak pinned at 100 (Figure 4).



4.2 Second page: a dashboard reporting 0.999 while dying

An early global ratio assessed memory quality, stable at 0.999 for ages, suspiciously healthy. We once took it for a grade. Root cause: the metric was a ratio over link types, not a count of real structure, so it never tracked the store's health. Concretely, the early metric was $\text{coherence} = (\text{total_links} - \text{contradict}) / \text{total_links}$; because contradict edges are rare across the whole store, the ratio pinned near 1 (0.9996, displayed as 0.999) and read green while the store degraded. The corrected metric counts true inter-standing-wave reinforcement density: $\text{coherence} = \log_{1p}(d) / \log_{1p}(3)$, where $d = 2 \times (\text{unique reinforce pairs}) / (\text{standing-wave node count})$, computed over the standing-wave pool alone. The standing-wave pool is the set of live cortex nodes that carry a completed vector and a weight above the activity gate; nodes below that gate stay in the raw layer and are not counted here. On the same store it read 0.101 (here $d \approx 0.15$). The divisor 3 is an empirical normalization constant and the score saturates at $d=3$; the metric is best read as ordinal rather than cardinal.

The most dangerous property of a structure-blind metric is not inaccuracy but the sicker the system, the greener the report stayed. A green dashboard does not mean the car is fine; it may only mean the dashboard is broken. Figure 5 shows the before/after. A metric should be a mirror—reflecting, never flattering; the 0.999 mirror reflected the lens, not the road.



4.3 Third page: load-bearing memories can score zero under popularity metrics


A full structural audit: 88% of raw items are islands: zero converging derivation. Operationally, we count an island as a raw utterance with $\sigma = 0$: no anchor-forming in-edges. The island share is a direct count over the raw-utterance cohort (user 67.8%, $n = 5,956$; AI 82.6%, $n = 3,173$; measured 2026-08). They are not garbage; they are roots. Two real examples (lightly redacted):

First, a self-description of decision style: *"His way of pricing risk: not by computing the worst case, but by looking for the minimum gain."* A dozen words, no derivation around them, zero on every counter. Its future value does not depend on being retrieved; it is the bedrock of how a person judges; some future decision will rest on it, invisible to any popularity metric.

Second, a passing remark about the roof of an old family house, unvisited for three months, zero derivation, zero counts, until a judgment of later consequence was formed and it turned out to be the sole basis. For those three months, retrieval popularity rated it as having no value.

This is the whole case for I2: root nodes are born cold; that is structure, not anomaly. A popularity-triggered cleanup can delete a root that later becomes a key judgment's sole basis; and after deletion, the dashboard gets even greener. The post-fix follow-up gives this page its coda: at that early audit the island share was 88%; by this writing it has fallen to 73% (user utterances 67.8%, AI utterances 82.6%; measured 2026-08, $n=9,129$). The fifteen-point drop is not the merit of cleanup but the merit of being left alone: roughly fourteen hundred raw

items (a 15-point drop on a cohort of about 9,100) have grown derivations with no one intervening. The bridge is not only on the blueprint; it is growing (Figure 6).



Case	Symptom	Hard evidence (measured)	Violates	Fix	Post-fix
§4.1	one field, three semantics	93× cross-type mean skew (449 vs 4.8); peak 1965 black hole	I3	split fields, per-semantics decay, cap [0,100]	93× → 2.2×; peak pinned at 100
§4.2	health reported in sickness	reported 0.999; structure-aware coherence 0.101	none — passes all of I0–I3	count real structure (inter-standing-wave reinforcement density)	now tracks the store
§4.3	popularity-driven deletion	audit 88% islands (now 73%), incl. the sole basis of a key judgment	I2 (write-time 1–10 scoring is an I0 violation)	lifecycle-triggered cascade metabolism	88% → 73%; ~1,400 items still deriving

Three pages, two failure axes. §4.1 and §4.3 are separation failures (I3, I2), each traceable to a well-meaning, convenient intervention. §4.2 belongs to a different axis: the invariants govern signal separation, not metric validity: the structure-blind metric passed all four invariants while measuring nothing. We record this as a distinct failure class: a checkup framework must also validate that a metric counts what it claims to count (see the seventh design law in §5).

5. Design Laws

1. **Salience is data, ranking is a service, metabolism is a lifecycle:** three signals, three mechanisms, physically separate fields.
2. **Counters never write back:** algorithmic processes may be recomputed from scratch; data they have altered cannot be restored.
3. **Measure salience in the correct scope:** counting edges across semantic domains counts fan-out.
4. **Metabolism cascades at the item level,** triggered by lifecycle; low frequency is not expiry.
5. **Polysemous fields are time bombs**—split early; every day late is another day of dirty values.
6. **Every importance metric must pass the Silence Test.**
7. **A metric must first be validated:** it must count what it claims to count. Structure-blind metrics pass every invariant while measuring nothing.

6. Tools: Two Audits You Can Run Today

The Silence Test. Suspend all retrieval behavior for T days (or replay history offline while intercepting all writes to persistent counters), then recompute your "importance/quality/salience" metrics. If the numbers do not move, you measured data; if they drift, you measured popularity: a health report contaminated by behavior. Cost: T days of patience, nearly zero.

The RAG Test (a blunter cut). Replace your entire "memory layer" with a good-enough search engine. Fix, in advance, a workload and an outcome criterion; if behavior shows no material difference on that pre-declared criterion, the layer was, by that measure, retrieval only—not memory. For memory to be memory, there must be evidence beyond the query: salience that grows while unretrieved, metabolism that happens while unobserved, memories that shape the model's behavior without being asked about.

Silence-Test Day: a five-step protocol (half a day).

1. Snapshot all current "importance/quality/salience" metric values;
2. Freeze all retrieval paths for T days, or replay history offline, intercepting every write to persistent counters;
3. Recompute the same metrics; compare item by item;
4. Every drift item is a behavior-contaminated item; mark its upstream write path (usually the $\partial s / \partial \text{access} > 0$ feedback source);
5. Publish the result, good or bad. A system passing I1 should show zero drift. One honest limit: a zero result isolates access-driven contamination only; contamination arising from other update paths must be controlled separately before independence is claimed. We followed this boundary up empirically on our own store: across the ~38k-node cortex, salience σ and the access counter are weakly but non-trivially correlated (Pearson 0.24, $n \approx 37,000$ live nodes), not zero. That residual comes from a shared structural cause rather than a write path: structurally central nodes are also more retrievable, and the correlation varies widely by node type (−0.17 to +0.50). The +0.50 within plan-type nodes is reported as it is: in that node type, salience and access are at present not separable, consistent with plan nodes being both structurally central and frequently recalled. We state this rather than trim it. The decisive evidence remains the formula: no access counter enters how σ is computed, and I1 constrains exactly that write path, not correlation.

We suggest maintainers schedule a Silence-Test Day and report results under this framework, favorable or not; what this field lacks is not new mechanisms but a first checkup of the old ones.

7. Discussion: Limitations, Validity Threats, and Why No Benchmark

Single system, single-user corpus, longitudinal; no cross-system benchmark (a concrete first trial is proposed below); salience is normalized by a store-level p99, so it drifts with store growth and requires renormalization for cross-store or cross-epoch comparison; it also depends on the association-edge structure and needs adaptation to migrate; the evidence comes from failure retrospectives and may carry survivor-narrative bias; we judge the direction unaffected, but state it plainly; the corpus is dominated by private conversation, so cross-domain generalization is unverified; "the life of one memory" is a worked example whose every step has corresponding production evidence (§4) but whose timeline is synthetic.

Why no cross-system benchmark? Three honest reasons. First, this paper is a single-system experience record: cross-system trials are future work, not a claim tested here. Second, mainstream benchmarks (e.g., LOCOMO [13]) measure question-answering accuracy, while our diagnostic target is signal structure; a system can score well on the benchmark while violating all four invariants (the 0.999 dashboard of §4.2 is exactly such "high-scoring sickness"). And the future work we propose is precisely to metricize the Silence Test itself: "the drift of the importance metric after T days of suspended retrieval" as a comparable cross-system memory-health metric (drift ≈ 0 passes). A concrete first trial: run the five-step Silence-Test Day on an open-source implementation (Memo or Zep) and report the drift; it requires no training compute at all, only honesty.

Data and artifact availability. All aggregate statistics (count distributions, island shares, before/after contrasts) appear as figures in this paper and in the public preprint (<https://doi.org/10.5281/zenodo.23068930>); the raw corpus will never be released for privacy; the five-step Silence-Test protocol is a public artifact, freely reusable and citable without permission.

Ethics note: the corpus originates from the author's own single-user conversational system; the quoted utterances are used with the author's consent.

8. Conclusion

Finally, the state of the system after the fix. Post-separation: cross-type distortion converged from $93\times$ to $2.2\times$, counter peaks pinned at 100; the quality metric became purely structural, counting real structure rather than edge ratios; the 88%-island raw utterances are alive to this day, many still deriving. Worth emphasizing: these three fixes introduced no new "clever algorithm" (no new model, no new scorer, no learned weights), only the three conflated responsibilities separated, so that salience and metabolism are decided by the data once more. That is all.

What we would rather leave with the reader is a less comfortable judgment: as long as the industry keeps intervening in LLM memory by hand, scoring in place of sedimentation, giving verdicts in place of metabolism, wiring human queries directly into memory past the middle medium, the cumulative effect is not better memory but a system that degrades under continued hand-tuning. RAG at least identifies itself as a retrieval engine; presenting one as a memory system while removing everything that happens beyond the query is a subtler and more misleading substitution.

Maintainers who have written such boosts, thresholds, or judge prompts are, in this sense, the executors of that degradation. The good news: the checkup sheet is in §6: five steps, half a day, and you can start today.

The three signals are now stored separately in our system. We invite maintainers to examine their own: many may find that the metric has been tracking retrieval behavior rather than memory all along.

References

- [1] Park, J.S., et al. *Generative Agents: Interactive Simulacra of Human Behavior*. arXiv:2304.03442, 2023.
- [2] Chhikara, P., et al. *Memo: Building Production-Ready AI Agents with Scalable Long-Term Memory*. ECAI 2025 / arXiv:2504.19413.
- [3] Memo Team. *Memory Eviction and Forgetting in AI Agents*. memo.ai/blog, 2026.
- [4] memoai/memo, GitHub Issue #5330: Memory Access Frequency Tracking and Lifetime-Based Cleanup.
- [5] Abdollahpouri, H., Burke, R., & Mobasher, B. *Controlling Popularity Bias in Learning-to-Rank Recommendation*. RecSys 2017.
- [6] Rasmussen, P.S., et al. *Zep: A Temporal Knowledge Graph Architecture for Agent Memory*. arXiv:2501.13956.
- [7] Orogat, A., & Mansour, E. *Is Agent Memory a Database? Rethinking Data Foundations for Long-Term AI Agent Memory*. arXiv:2605.26252, 2026.
- [8] Roediger, H.L., & Karpicke, J.D. *Test-Enhanced Learning: Taking Memory Tests Improves Long-Term Retention*. Psychological Science, 2006.
- [9] Rasch, B., & Born, J. *About Sleep's Role in Memory*. Physiological Reviews, 2013.
- [10] Lewis, P., et al. *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. NeurIPS 2020.

- [11] Zhong, W., et al. *MemoryBank: Enhancing Large Language Models with Long-Term Memory*. AAAI 2024.
 - [12] Packer, C., et al. *MemGPT: Towards LLMs as Operating Systems*. arXiv:2310.08560, 2023.
 - [13] Maharana, A., et al. *Evaluating Very Long-Term Conversational Memory of LLM Agents*. ACL 2024.
-